

The University of New South Wales

COMP9315 DBMS Implementation

Final Exam 14s2

[\[Instructions\]](#) [\[Notes\]](#) [\[PostgreSQL\]](#) [\[C\]](#)
[\[Q1\]](#) [\[Q2\]](#) [\[Q3\]](#) [\[Q4\]](#) [\[Q5\]](#) [\[Q6\]](#) [\[Q7\]](#) [\[Q8\]](#)

Question 1 (20 marks)

For this question, you will write functions to implement two of the major operations of an *iterator*. This iterator, called `scan`, fetches all of the tuples from a file, getting the tuples in the order that they are stored in the file.

The data files used by this iterator have the following properties:

- a **file** is a sequence of N pages numbered $0..N-1$
- a **page** contains a header and a sequence of M tuples, numbered $0..M-1$
- a **tuple** is a sequence of characters, terminated by `'\0'` (i.e. a C string)

A page is a 1024-byte block, structured as follows (see `Page.h` for details)

- a header consisting of a 4-byte counter followed by an array of 63 offsets
- an array of bytes, containing tuple strings

The following diagram shows the detailed structure of each page:

nTuples	tup0offset	tup1offset
tup2offset	tup3offset
etc etc etc (more offsets)		
tup0:"a1,b1,c1"		tup1:"a2,b2,c2"
tup2:"a3,b3,c3"	tup3:"a4,b4,c4"	
etc etc etc (more tuples)		

The iterator has a main program which works as a tuple scanner. The `main()` function for the scanner takes a file name as a command-line argument, opens the file, starts the iterator and then fetches and prints all of the tuples in the file. The scanner also accesses the `scan` data structure to print out the current page and current tuple number in the page of the fetched tuple. There are examples of the output from `scanner` in the directory `q1/tests`, in the files with names like `t3.expected`.

The following files are available for the implementation and can be found in the `q1` directory in your exam environment:

- `Makefile` ... to control the compilation of the scanner
- `main.c` ... the main program of the scanner
- `Page.h` ... data structures for pages
- `Page.c` ... operations on pages
- `Scan.h` ... data structures for the iterator
- `Scan.c` ... operations on the iterator (incomplete)

All of these files are complete and fully-functional except for `Scan.c`, which you are required to complete.

If you run the `make` command, it will produce an executable file called `scanner`, which you could use to try to read the tuples from the test files using e.g.

```
$ ./scanner tests/t3
```

At this stage, the only output you'll get is "Cannot start scan". Once you have correctly implemented the iterator, you'll see the tuples instead.

Note that the data files are contained in a sub-directory called `tests`, along with the expected output for when a correct `scanner` is run on them. Some hopefully useful notes about the data files:

- `t0` is a completely empty data file (0 bytes)
- `t1` contains a single page, but the page contains zero tuples
- `t2` contains a single page, with 15 tuples
- other data files contain a mixture of empty and non-empty pages
- offsets are relative to the start of the tuples, *not* the start of the Page
- each page is zero-filled when it is initially created (before tuples are added)

To help you check whether your program is working correctly, there is a script called `check` which will run your program against all of the tests and report the results, and add some output files to the `tests` directory which might help you to debug your code. Since `check` produces a lot of output (until your program is working), it might be useful to run it like:

```
$ check | less
```

Your task is to complete the two functions `startScan()` and `nextTuple()` in `Scan.c`. You must preserve the function interfaces and you must use the supplied `Scan` structure. You are not allowed to modify any of the other files (except to add debugging code), and you cannot submit any file apart from `Scan.c`. It requires around 30 lines of code to solve this problem; partial marks are available if you complete some of the code.

Some hints on how to approach this problem:

- take a quick look at the `main.c` file to see what it does and how it uses the `startScan()` and `nextTuple()` functions
- read the data structure definitions in `Scan.h` and `Page.h`
- start by implementing `startScan()`
- think carefully about the way `nextTuple()` moves from tuple to tuple and page to page
- use `gdb`, if you know how; otherwise add plenty of `printf`'s for debugging

The directory `tests` contains data files (called `ti`) and an expected output file (called `ti.out`) from running `./scanner` on each data file. You should look at the expected output files to see what a correct program should produce. The data files are in binary format, so can't be usefully viewed with `cat` or `less` or a text editor. If you want to look at them, try the `od` command (e.g. `od -c tests/t3`).

Once your program passes the `check` tests with no errors, you can submit it.

Instructions:

- Type your answer to this question into the file called `Scan.c`
- Submit via: **submit q1**

(The `submit` command knows to collect the `Scan.c` file)

Hint: you may want to leave this question until you have completed all of the other questions.

End of Question