

B-trees

- B-Trees
- B-Tree Depth
- Selection with B-Trees
- Insertion into B-Trees
- Example: B-tree Insertion
- B-Tree Insertion Cost
- B-trees in PostgreSQL

❖ B-Trees

^ >>

B-trees are multi-way search trees with the properties:

- they are updated so as to remain balanced
- each node has at least $(n-1)/2$ entries in it
- each tree node occupies an entire disk page

B-tree insertion and deletion methods

- are moderately complicated to describe
- can be implemented very efficiently

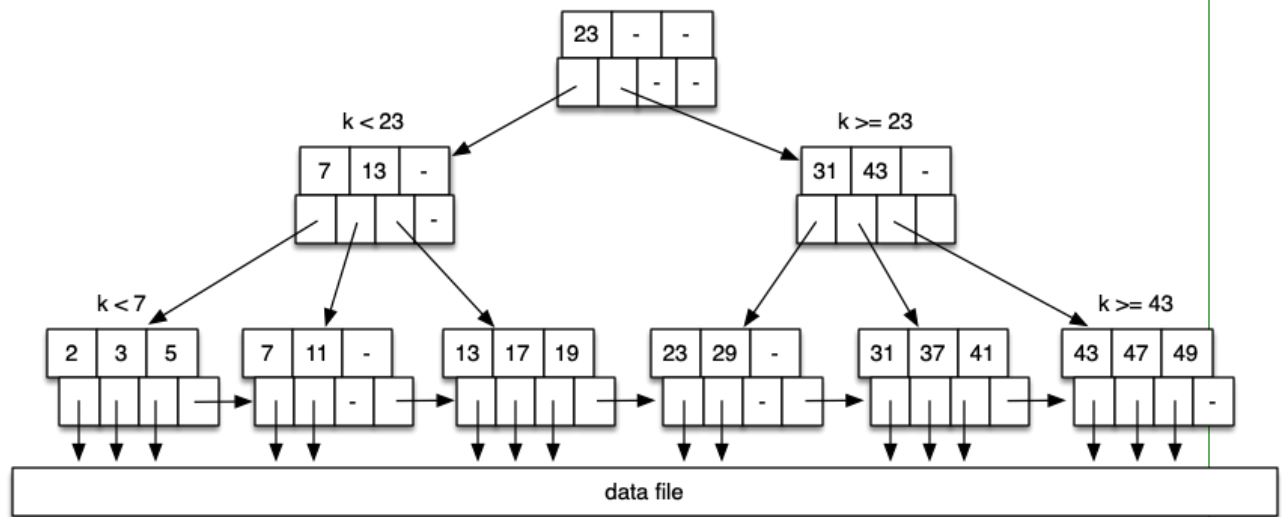
Advantages of B-trees over general multi-way search trees:

- better storage utilisation (around 2/3 full)
- better worst case performance (shallower)

COMP9315 21T1 ◇ B-trees ◇ [1/15]

❖ B-Trees (cont)

Example B-tree (depth=3, $n=3$) (actually B+ tree)



(Note: in DBs, nodes are pages \Rightarrow large branching factor, e.g. $n=500$)

❖ B-Tree Depth

Depth depends on effective branching factor (i.e. how full nodes are).

Simulation studies show typical B-tree nodes are 69% full.

Gives load $L_i = 0.69 \times c_i$ and depth of tree $\sim \text{ceil}(\log_{L_i} r)$.

Example: $c_i = 128$, $L_i = 88$

Level	#nodes	#keys
root	1	88
1	89	7832
2	7921	697048
3	704969	62037272

Note: c_i is generally larger than 128 for a real B-tree.

❖ Selection with B-Trees

For *one* queries:

```
Node find(k, tree) {
    return search(k, root_of(tree))
}
Node search(k, node) {
    // get the page of the node
    if (is_leaf(node)) return node
    keys = array of nk key values in node
    pages = array of nk+1 ptrs to child nodes
    if (k <= keys[0])
        return search(k, pages[0])
    else if (keys[i] < k <= keys[i+1])
        return search(k, pages[i+1])
    else if (k > keys[nk-1])
        return search(k, pages[nk])
}
```

❖ Selection with B-Trees (cont)

Simplified description of search ...

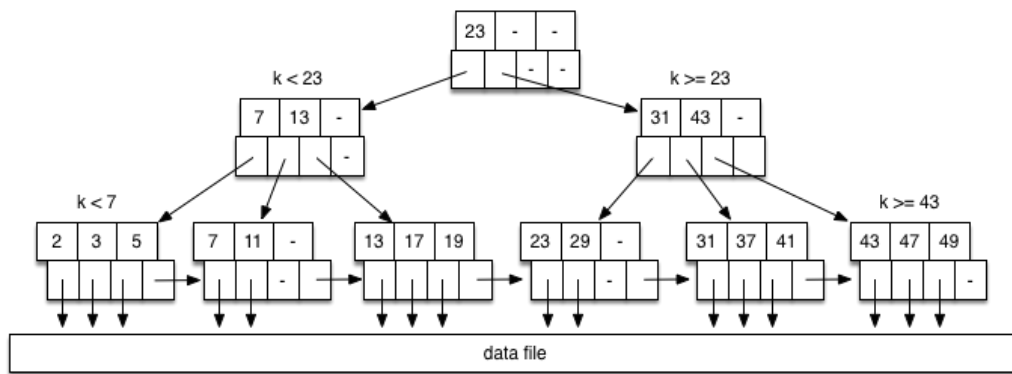
N = B-tree root node

while (N is not a leaf node) $N = \text{scanToFindChild}(N, K)$

$\text{tid} = \text{scanToFindEntry}(N, K)$

access tuple T using tid

$$\text{Cost}_{\text{one}} = (D + 1)_r$$



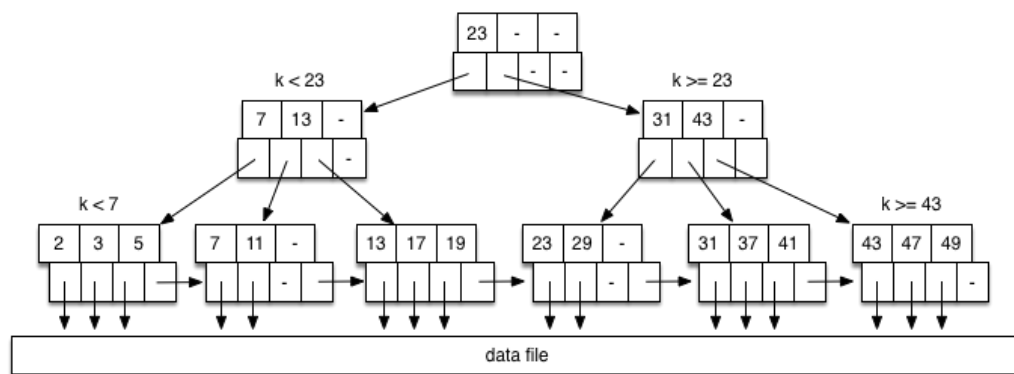
❖ Selection with B-Trees (cont)

For *range* queries (assume sorted on index attribute):

```

search index to find leaf node for Lo
for each leaf node entry until Hi found
    add pageOf(tid) to Pages to be scanned
scan Pages looking for matching tuples
  
```

$$Cost_{range} = (D + b_i + b_q)_r$$



❖ Insertion into B-Trees


Overview of the method:

1. find leaf node and position in node where new key belongs
2. if node is not full, insert entry into appropriate position
3. if node is full ...
 - promote middle element to parent
 - split node into two half-full nodes ($< \text{middle}$, $> \text{middle}$)
 - insert new key into appropriate half-full node
4. if parent full, split and promote upwards
5. if reach root, and root is full, make new root upwards

Note: if duplicates not allowed and key exists, may stop after step 1.

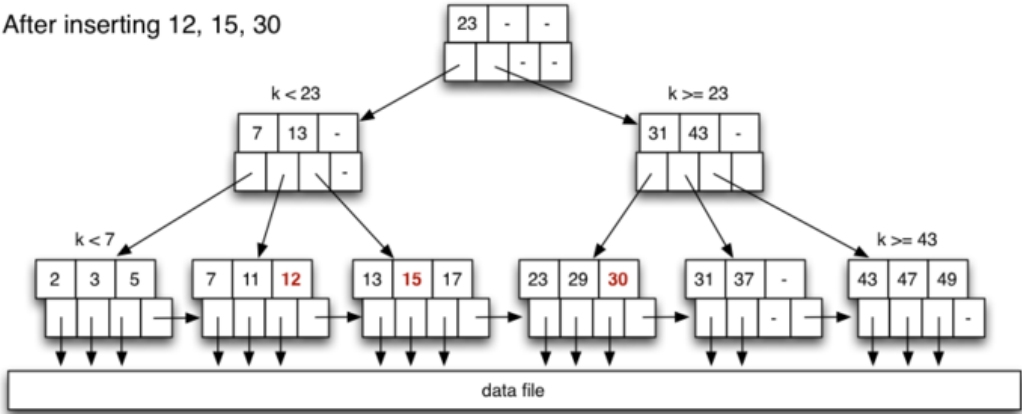
❖ Example: B-tree Insertion

Starting from this tree:





insert the following keys in the given order **12 15 30 10**

❖ Example: B-tree Insertion (cont)



❖ Example: B-tree Insertion (cont)



❖ B-Tree Insertion Cost

Insertion cost = $Cost_{treeSearch} + Cost_{treeInsert} + Cost_{dataInsert}$

Best case: write one page (most of time)

- traverse from root to leaf
- read/write data page, write updated leaf

$$Cost_{insert} = D_r + 1_w + 1_r + 1_w$$

Common case: 3 node writes (rearrange 2 leaves + parent)

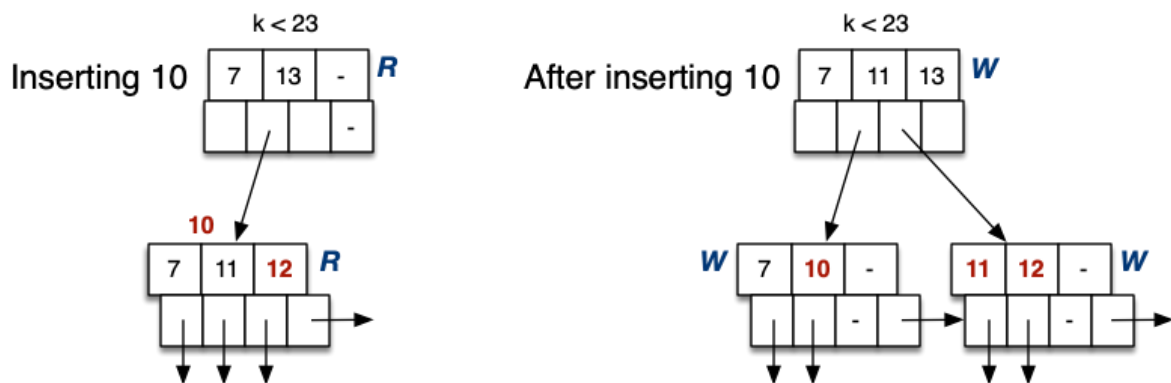
- traverse from root to leaf, holding nodes in buffer
- read/write data page
- update/write leaf, parent and sibling

$$Cost_{insert} = D_r + 3_w + 1_r + 1_w$$

❖ B-Tree Insertion Cost (cont)

Worst case: propagate to root $Cost_{insert} = D_r + D \cdot 3W + 1_r + 1_w$

- traverse from root to leaf
- read/write data page
- update/write leaf, parent and sibling
- repeat previous step $D-1$ times



❖ B-trees in PostgreSQL

PostgreSQL implements \cong Lehman/Yao-style B-trees

- variant that works effectively in high-concurrency environments.

B-tree implementation: **backend/access/nbtree**

- **README** ... comprehensive description of methods
- **nbtree.c** ... interface functions (for iterators)
- **nbtsearch.c** ... traverse index to find key value
- **nbtinsert.c** ... add new entry to B-tree index

Notes:

- stores all instances of equal keys (dense index)
- avoids splitting by scanning right if key = max(key) in page
- common insert case: new key is max(key) overall; handled efficiently

❖ B-trees in PostgreSQL (cont)

Changes for PostgreSQL v12

- indexes smaller
 - for composite keys, only store first attribute
 - index entries are smaller, so c_i larger, so tree shallower
- include TID in index key
 - duplicate index entries are stored in "table order"
 - makes scanning table files to collect results more efficient

To explore indexes in more detail:

- `\di+ IndexName`
- `select * from bt_page_items (IndexName, BlockNo)`

❖ B-trees in PostgreSQL (cont)

Interface functions for B-trees

```
// build Btree index on relation
Datum btbuild(rel,index,...)
// insert index entry into Btree
Datum btinsert(rel,key,tupleid,index,...)
// start scan on Btree index
Datum btbeginscan(rel,key,scandesc,...)
// get next tuple in a scan
Datum btgettupple(scandesc,scandir,...)
// close down a scan
Datum btendscan(scandesc)
```


Produced: 24 Mar 2021