

Implementing Join

- Join
- Join Example
- Implementing Join
- Join Summary
- Join in PostgreSQL

❖ Join

DBMSs are engines to **store**, **combine** and **filter** information.

Join (\bowtie) is the primary means of **combining** information.

Join is important and potentially expensive

Most common join condition: equijoin, e.g. **(R.pk = S.fk)**

Join varieties (natural, inner, outer, semi, anti) all behave similarly.

We consider three strategies for implementing join

- **nested loop** ... simple, widely applicable, inefficient without buffering
- **sort-merge** ... works best if tables are sorted on join attributes
- **hash-based** ... requires good hash function and sufficient buffering

❖ Join Example

Consider a university database with the schema:

```
create table Student(  
    id      integer primary key,  
    name    text, ...  
);  
create table Enrolled(  
    stude   integer references Student(id),  
    subj    text references Subject(code), ...  
);  
create table Subject(  
    code    text primary key,  
    title   text, ...  
);
```

We use this example for each join implementation, by way of comparison

❖ Join Example (cont)

Goal: *List names of students in all subjects, arranged by subject.*

SQL query to provide this information:

```
select E.subj, S.name
from   Student S
       join Enrolled E on (S.id = E.stude)
order  by E.subj, S.name;
```

And its relational algebra equivalent:

$Sort[subj] (Project[subj,name] (Join[id=stude] (Student, Enrolled)))$

To simplify formulae, we denote **Student** by S and **Enrolled** by E

❖ Join Example (cont)

Some database statistics:

Sym	Meaning	Value
r_S	# student records	20,000
r_E	# enrollment records	80,000
c_S	Student records/page	20
c_E	Enrolled records/page	40
b_S	# data pages in Student	1,000
b_E	# data pages in Enrolled	2,000

Also, in cost analyses later, N = number of memory buffers.

❖ Join Example (cont)

Relation statistics for **Out** = *Student* ⋈ *Enrolled*

Sym	Meaning	Value
r_{Out}	# tuples in result	80,000
C_{Out}	result records/page	80
b_{Out}	# data pages in result	1,000

Notes:

- r_{Out} ... one result tuple for each **Enrolled** tuple
- C_{Out} ... result tuples have only **subj** and **name**
- in analyses, ignore cost of writing result ... same in all methods

❖ Implementing Join

A naive join implementation strategy

```
for each tuple  $T_S$  in Students {  
  for each tuple  $T_E$  in Enrolled {  
    if (testJoinCondition( $C, T_S, T_E$ )) {  
       $T1 = \text{concat}(T_S, T_E)$   
       $T2 = \text{project}([\text{subj}, \text{name}], T1)$   
       $\text{ResultSet} = \text{ResultSet} \cup \{T2\}$   
    }  
  }  
}
```

Problems:

- join condition is tested $r_E \cdot r_S = 16 \times 10^8$ times
- tuples scanned = $r_S + r_S \cdot r_E = 20000 + 20000 \cdot 80000 = 1600020000$

❖ Implementing Join (cont)

An alternative naive join implementation strategy

```
for each tuple  $T_E$  in Enrolled {  
  for each tuple  $T_S$  in Students {  
    if (testJoinCondition( $C, T_S, T_E$ )) {  
       $T1 = \text{concat}(T_S, T_E)$   
       $T2 = \text{project}([\text{subj}, \text{name}], T1)$   
       $\text{ResultSet} = \text{ResultSet} \cup \{T2\}$   
    }  
  }  
}
```

Relatively minor performance difference ...

- tuples scanned = $r_E + r_E \cdot r_S = 80000 + 80000 \cdot 20000 = 1600080000$

Terminology: relation in outer loop is **outer**; other relation is **inner**

❖ Join Summary

None of **nested-loop**/**sort-merge**/**hash** join is superior in some overall sense.

Which strategy is best for a given query depends on:

- sizes of relations being joined, size of buffer pool
- any indexing on relations, whether relations are sorted
- which attributes and operations are used in the query
- number of tuples in S matching each tuple in R
- distribution of data values (uniform, skew, ...)

Given query Q , choosing the "best" join strategy is critical;
the cost difference between best and worst case can be very large.

E.g. $Join[id=stude](Student, Enrolled)$: 3,000 ... 2,000,000 page reads

❖ Join in PostgreSQL

Join implementations are under:

src/backend/executor

PostgreSQL supports the three join methods that we discuss:

- nested loop join (**nodeNestloop.c**)
- sort-merge join (**nodeMergejoin.c**)
- hash join (**nodeHashjoin.c**) (hybrid hash join)

The query optimiser chooses appropriate join, by considering

- physical characteristics of tables being joined
- estimated selectivity (likely number of result tuples)

Produced: 22 Mar 2021