

PostgreSQL Tuples

- PostgreSQL Tuples
- PostgreSQL Attribute Values

❖ PostgreSQL Tuples

Definitions: `include/postgres.h`, `include/access/*tup*.h`

Functions: `backend/access/common/*tup*.c` e.g.

- `HeapTuple heap_form_tuple(desc, values[], isnull[])`
- `heap_deform_tuple(tuple, desc, values[], isnull[])`

PostgreSQL implements tuples via:

- a contiguous chunk of memory
- starting with a header giving e.g. #fields, nulls
- followed by data values (as a sequence of **Datum**)

❖ PostgreSQL Tuples (cont)

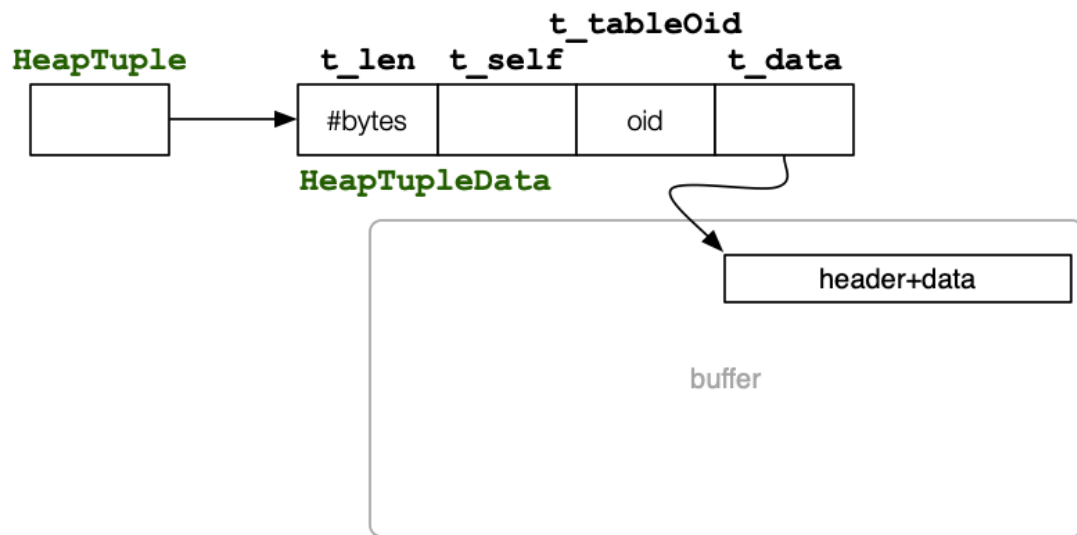
HeapTupleData contains information about a stored tuple

```
typedef HeapTupleData *HeapTuple;  
  
typedef struct HeapTupleData  
{  
    uint32          t_len;    // length of *t_data  
    ItemPointerData t_self;   // SelfItemPointer  
    Oid             t_tableOid; // table the tuple came from  
    HeapTupleHeader t_data;   // -> tuple header and data  
} HeapTupleData;
```

HeapTupleHeader is a pointer to a location in a buffer

❖ PostgreSQL Tuples (cont)

Structure of **HeapTuple**:



❖ PostgreSQL Tuples (cont)

PostgreSQL stores each record as tuple header, followed by data:

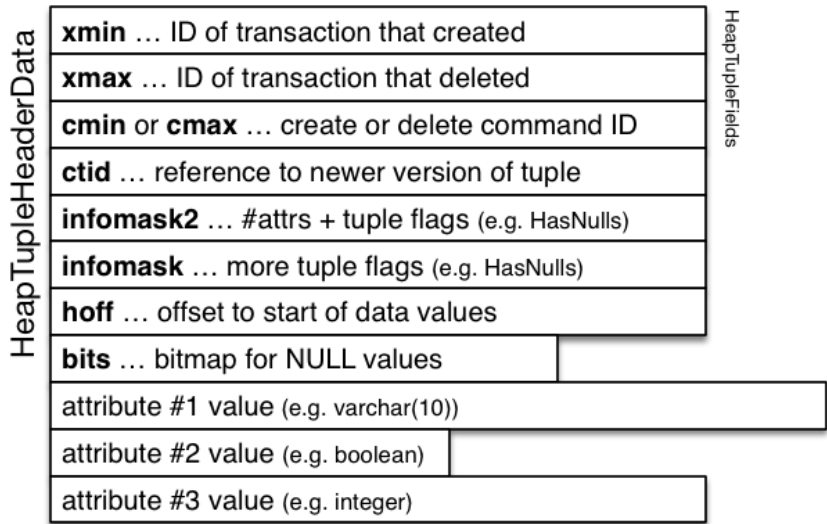
```
typedef HeapTupleHeaderData *HeapTupleHeader;  
  
typedef struct HeapTupleHeaderData // simplified  
{  
    HeapTupleFields t_heap;  
    ItemPointerData t_ctid;          // TID of newer version  
    uint16          t_infomask2;    // #attributes + flags  
    uint16          t_infomask;     // flags e.g. has_null  
    uint8           t_hoff;         // sizeof header incl. t_bits  
    // above is fixed size (23 bytes) for all heap tuples  
    bits8           t_bits[1];     // bitmap of NULLs, var.len.  
    // OID goes here if HEAP_HASOID is set in t_infomask  
    // actual data follows at end of struct  
} HeapTupleHeaderData;
```

COMP9315 21T1 ◇ PostgreSQL Tuples ◇ [4/12]

<< ^ >>

❖ PostgreSQL Tuples (cont)

Tuple structure:



❖ PostgreSQL Tuples (cont)

Some of the bits in **t_infomask**..

```
#define HEAP_HASNULL      0x0001
/* has null attribute(s) */
#define HEAP_HASVARWIDTH  0x0002
/* has variable-width attribute(s) */
#define HEAP_HASEXTERNAL  0x0004
/* has external stored attribute(s) */
#define HEAP_HASOID_OLD   0x0008
/* has an object-id field */
```

Location of **NULLs** is stored in **t_bits[]** array

❖ PostgreSQL Tuples (cont)

Tuple-related data types: (cont)

```
typedef struct HeapTupleFields // simplified
{
    TransactionId t_xmin; // inserting xact ID
    TransactionId t_xmax; // deleting or locking xact ID
    union {
        CommandId t_cid; // inserting or deleting command ID
        TransactionId t_xvac; // old-style VACUUM FULL xact ID
    } t_field3;
} HeapTupleFields;
```

Note that not all system fields from stored tuple appear

- **oid** is stored after the tuple header, if used
- both **xmin/xmax** are stored, but only one of **cmin/cmax**

❖ PostgreSQL Tuples (cont)

Tuple-related data types: (cont)

```
// TupleDesc: schema-related information for HeapTuples

typedef struct tupleDesc
{
    int          natts;          // # attributes in tuple
    Oid          tdtypeid;       // composite type ID for tuple type
    int32        tdtypmod;       // typmod for tuple type
    bool         tdhasoid;       // does tuple have oid attribute?
    int          tdrefcount;     // reference count (-1 if not counting)
    TupleConstr *constr;         // constraints, or NULL if none
    FormData_pg_attribute attrs[];
    // attrs[N] is a pointer to description of attribute N+1
} *TupleDesc;
```

COMP9315 21T1 ◇ PostgreSQL Tuples ◇ [8/12]

❖ PostgreSQL Tuples (cont)

Tuple-related data types: (cont)

```
// FormData_pg_attribute:  
// schema-related information for one attribute  
  
typedef struct FormData_pg_attribute  
{  
    Oid          attrelid;      // OID of reln containing attr  
    NameData attname;          // name of attribute  
    Oid          atttypid;      // OID of attribute's data type  
    int16        attlen;        // attribute length  
    int32        attndims;      // # dimensions if array type  
    bool         attnotnull;    // can attribute have NULL value  
    .....                // and many other fields  
} FormData_pg_attribute;
```

For details, see **include/catalog/pg_attribute.h**

❖ PostgreSQL Attribute Values

Attribute values in PostgreSQL tuples are packaged as **Datum**s

```
// representation of a data value  
typedef uintptr_t Datum;
```

The actual data value:

- may be stored in the **Datum** (e.g. **int**)
- may have a header with length (for varlen attributes)
- may be stored in a TOAST file (if large value)

❖ PostgreSQL Attribute Values (cont)

Attribute values can be extracted as **Datum** from **HeapTuples**

```
Datum heap_getattr(  
    HeapTuple tup,          // tuple (in memory)  
    int attnum,             // which attribute  
    TupleDesc tupDesc,     // field descriptors  
    bool *isnull           // flag to record NULL  
)
```

isnull is set to true if value of field is **NULL**

attnum can be negative ... to access system attributes (e.g. OID)

For details, see **include/access/htup_details.h**

❖ PostgreSQL Attribute Values (cont)

Values of **Datum** objects can be manipulated via e.g.

```
// DatumGetBool:  
//   Returns boolean value of a Datum.  
  
#define DatumGetBool(X) ((bool) ((X) != 0))  
  
// BoolGetDatum:  
//   Returns Datum representation for a boolean.  
  
#define BoolGetDatum(X) ((Datum) ((X) ? 1 : 0))
```

For details, see **include/postgres.h**

Produced: 28 Feb 2021