

Signature-based Indexing

- Indexing with Signatures
- Signatures
- Generating Codewords
- Superimposed Codewords (SIMC)
- Concatenated Codewords (CATC)
- Queries using Signatures
- False Matches
- SIMC vs CATC



❖ Indexing with Signatures

Signature-based indexing:

- designed for *pmr* queries (conjunction of equalities)
- does not try to achieve better than $O(n)$ performance
- attempts to provide an "efficient" linear scan

Each tuple is associated with a **signature**

- a compact (lossy) descriptor for the tuple
- formed by combining information from multiple attributes
- stored in a signature file, parallel to data file

Instead of scanning/testing tuples, do pre-filtering via signatures.

❖ Indexing with Signatures (cont)

File organisation for signature indexing (two files)

Signature File

[0]	
[1]	
[2]	
[3]	
[4]	
[5]	
[6]	
[7]	
[8]	
...	

Data File

[0]	
[1]	
[2]	
[3]	
[4]	
[5]	
[6]	
[7]	
[8]	
...	

One signature slot per tuple slot; unused signature slots are zeroed.

Signatures do not determine record placement \Rightarrow can use with other indexing.

❖ Signatures

A **signature** "summarises" the data from one tuple

A tuple consists of n attribute values $A_1 .. A_n$

A **codeword** $cw(A_i)$ is

- a bit-string, m bits long, where k bits are set to 1 ($k \ll m$)
- derived from the value of a single attribute A_i

A **tuple descriptor** (signature) is built

- by combining $cw(A_i)$, $i=1..n$
- aim to have roughly half of the bits set to 1

Two strategies for building signatures: overlay, concatenate

❖ Generating Codewords

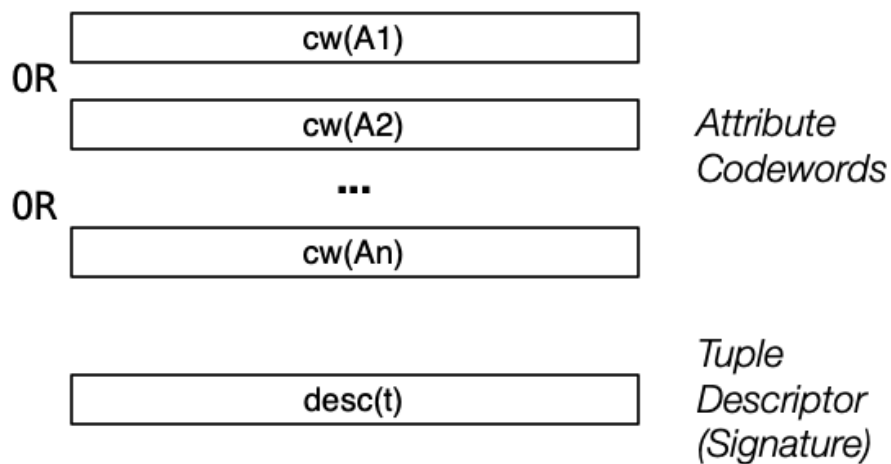
Generating a k -in- m codeword for attribute A_i

```
bits codeword(char *attr_value, int m, int k)
{
    int  nbits = 0;    // count of set bits
    bits cword = 0;    // assuming m <= 32 bits
    srand(hash(attr_value));
    while (nbits < k) {
        int i = random() % m;
        if (((1 << i) & cword) == 0) {
            cword |= (1 << i);
            nbits++;
        }
    }
    return cword;    // m-bits with k 1-bits and m-k 0-bits
}
```

❖ Superimposed Codewords (SIMC)

In a superimposed codewords (simc) indexing scheme

- tuple descriptor formed by *overlaying* attribute codewords (bitwise-OR)



❖ Superimposed Codewords (SIMC) (cont)

A SIMC tuple descriptor $desc(t)$ is

- a bit-string, m bits long, where $j \leq nk$ bits are set to 1
- $desc(t) = cw(A_1) \text{ OR } cw(A_2) \text{ OR } \dots \text{ OR } cw(A_n)$

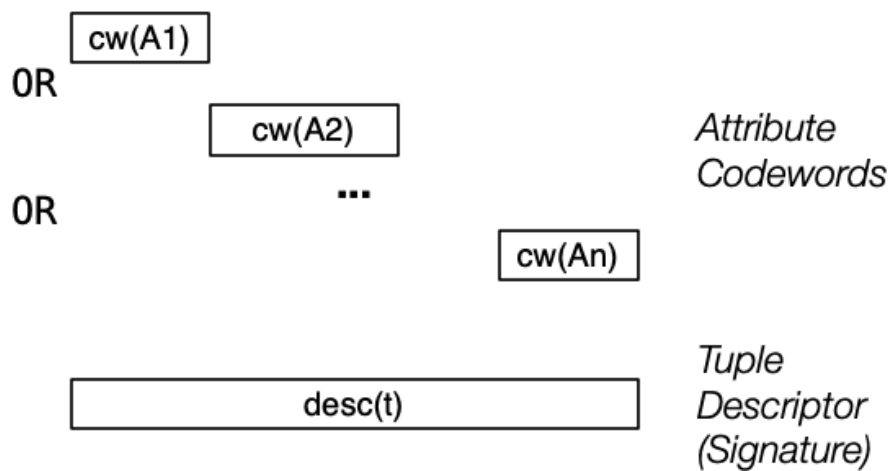
Method (assuming all n attributes are used in descriptor):

```
Bits desc = 0
for (i = 1; i <= n; i++) {
    bits cw = codeword(A[i], m, k)
    desc = desc | cw
}
```

❖ Concatenated Codewords (CATC)

In a concatenated codewords (catc) indexing schema

- tuple descriptor formed by *concatenating* attribute codewords



❖ Concatenated Codewords (CATC) (cont)

A CATC tuple descriptor $desc(t)$ is

- a bit-string, m bits long, where $j = nk$ bits are set to 1
- $desc(t) = cw(A_1) + cw(A_2) + \dots + cw(A_n)$ (+ is concatenation)

Each codeword is $p = m/n$ bits long, with $k = p/2$ bits set to 1

Method (assuming all n attributes are used in descriptor):

```
Bits desc = 0 ;  int p = m/n
for (i = 1; i <= n; i++) {
    bits cw = codeword(A[i], p, k)
    desc = desc | (cw << p*(n-i))
}
```

❖ Queries using Signatures

To answer query q with a signature-based index

- first generate a **query descriptor** $desc(q)$
- then scan the signature file using the query descriptor
- if sig_i matches $desc(q)$, then tuple i may be a match

$desc(q)$ is formed from codewords of known attributes.

Effectively, any unknown attribute A_i has $cw(A_i) = 0$

E.g. for SIMC $(a, ?, c, ?, e) = cw(a) \text{ OR } 0 \text{ OR } cw(c) \text{ OR } 0 \text{ OR } cw(e)$

❖ Queries using Signatures (cont)

Once we have a query descriptor, we search the signature file:

```
pagesToCheck = {}  
// scan r descriptors  
for each descriptor D[i] in signature file {  
    if (matches(D[i], desc(q))) {  
        pid = pageOf(tupleID(i))  
        pagesToCheck = pagesToCheck U pid  
    }  
}  
// scan  $b_q + \delta$  data pages  
for each pid in pagesToCheck {  
    Buf = getPage(dataFile, pid)  
    check tuples in Buf for answers  
}
```

❖ False Matches

Both SIMC and CATC can produce **false matches**

- **matches**(**D**[**i**] , **desc**(**q**)) is true, but **Tup**[**i**] is not a solution for **q**

Why does this happen?

- signatures are based on hashing, and it is possible that
hash(key₁) == hash(key₂) even though **key₁ != key₂**
- for SIMC, overlaying could also produce "unfortunate" bit-combinations

To mitigate this, need to choose "good" *m* and *k*

❖ SIMC vs CATC

Both build m -bit wide signatures, with $\sim 1/2$ bits set to 1

Both have codewords with $\sim m/2n$ bits set to 1

CATC: codewords are $m/n = p$ -bits wide

- shorter codewords \rightarrow more hash collisions

SIMC: codewords are also m -bits wide

- longer codewords \Rightarrow less hash collisions, but also has overlay collisions

CATC has option of different length codeword p_i for each A_i
($\sum p_i = m$)

Produced: 20 Mar 2021