

Page Internals

- Pages
- Page Formats
- Page Formats
- Storage Utilisation
- Overflows

❖ Pages

Database applications view data as:

- a collection of records (tuples)
- records can be accessed via a **TupleId/RecordId/RID**
- **TupleId = (PageID + TupIndex)**


The disk and buffer manager provide the following view:

- data is a sequence of fixed-size pages (aka "blocks")
- pages can be (random) accessed via a **PageID**
- each page contains zero or more tuple values

Page format = how space/tuples are organised within a page

❖ Pages (cont)

Data files consist of pages containing tuples:



*r tuples contained in b pages
each page can hold up to c tuples*

Each data file (in PostgreSQL) is related to one table.

❖ Page Formats

Ultimately, a **Page** is simply an array of bytes (**byte[]**).

We want to interpret/manipulate it as a collection of **Records** (tuples).

Tuples are addressed by a record ID (**rid = (PageId, TupIndex)**)

Typical operations on **Pages**:

- **request_page(pid)** ... get page via its **PageId**
- **get_record(rid)** ... get record via its **TupleId**
- **rid = insert_record(pid, rec)** ... add new record
- **update_record(rid, rec)** ... update value of record
- **delete_record(rid)** ... remove record from page

❖ Page Formats (cont)

Page format = tuples + data structures allowing tuples to be found

Characteristics of **Page** formats:


- record size variability (fixed, variable)
- how free space within **Page** is managed
- whether some data is stored outside **Page**
 - does **Page** have an associated overflow chain?
 - are large data values stored elsewhere? (e.g. TOAST)
 - can one tuple span multiple **Pages**?

Implementation of **Page** operations critically depends on format.

❖ Page Formats (cont)

For fixed-length records, use **record slots**.

- **insert**: place new record in first available slot
- **delete**: two possibilities for handling free record slots:



❖ Page Formats

For variable-length records, must use **slot directory**.

Possibilities for handling free-space within block:

- compacted (one region of free space)
- fragmented (distributed free space)

In practice, a combination is useful:


- normally fragmented (cheap to maintain)
- compacted when needed (e.g. record won't fit)

Important aspect of using slot directory

- location of tuple within page can change, tuple index does not change

❖ Page Formats (cont)


Compacted free space:



Note: "pointers" are implemented as word offsets within block.


❖ Page Formats (cont)

Fragmented free space:




❖ Page Formats (cont)

Initial page state (compacted free space) ...



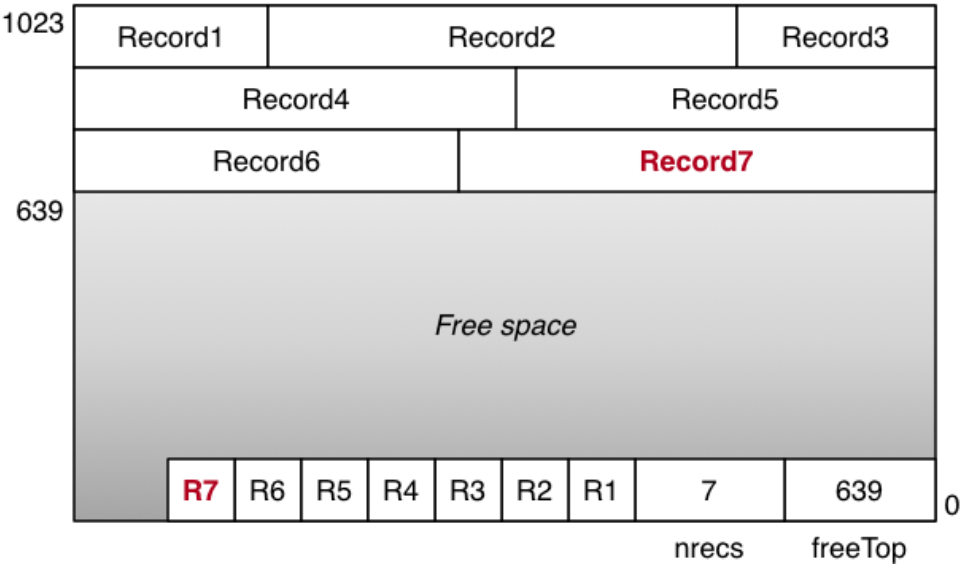
❖ Page Formats (cont)

Before inserting record 7 (compacted free space) ...



❖ Page Formats (cont)


After inserting record 7 (80 bytes) ...



<< ^ >>

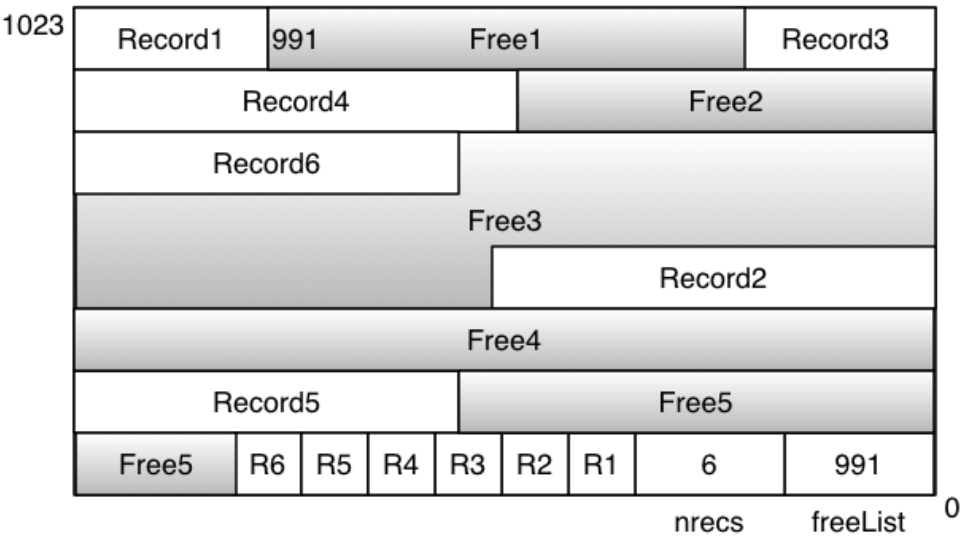
❖ Page Formats (cont)

Initial page state (fragmented free space) ...



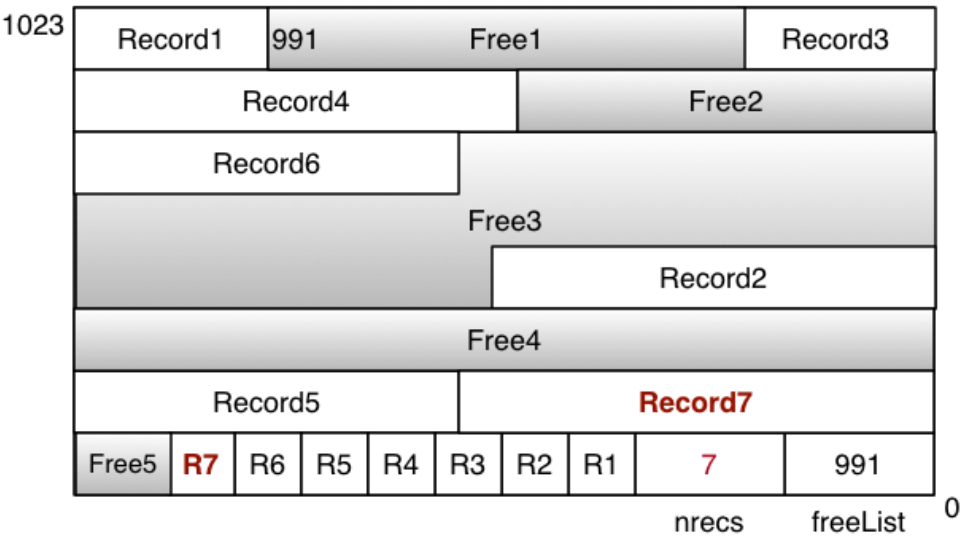
❖ Page Formats (cont)

Before inserting record 7 (fragmented free space) ...



❖ Page Formats (cont)

After inserting record 7 (80 bytes) ...



❖ Storage Utilisation

How many records can fit in a page? (denoted C = capacity)

Depends on:

- page size ... typical values: 1KB, 2KB, 4KB, 8KB
- record size ... typical values: 64B, 200B, app-dependent
- page header data ... typically: 4B - 32B
- slot directory ... depends on how many records

We typically consider *average* record size (R)

Given C , $HeaderSize + C * SlotSize + C * R \leq PageSize$

❖ Overflows

Sometimes, it may not be possible to insert a record into a page:

1. no free-space fragment large enough
2. overall free-space is not large enough
3. the record is larger than the page
4. no more free directory slots in page

For case (1), can first try to compact free-space within the page.

If still insufficient space, we need an alternative solution ...

❖ Overflows (cont)

File organisation determines how cases (2)..(4) are handled.

If records may be inserted anywhere that there is free space

- cases (2) and (4) can be handled by making a new page
- case (3) requires either spanned records or "overflow file"


If file organisation determines record placement (e.g. hashed file)

- cases (2) and (4) require an "overflow page"
- case (3) requires an "overflow file"


With overflow pages, *rid* structure may need modifying (*rel,page,ovfl,rec*)

❖ Overflows (cont)

Overflow files for very large records and BLOBs:



Record-based handling of overflows:



We discuss overflow pages in more detail when covering Hash Files.

Produced: 23 Feb 2021