

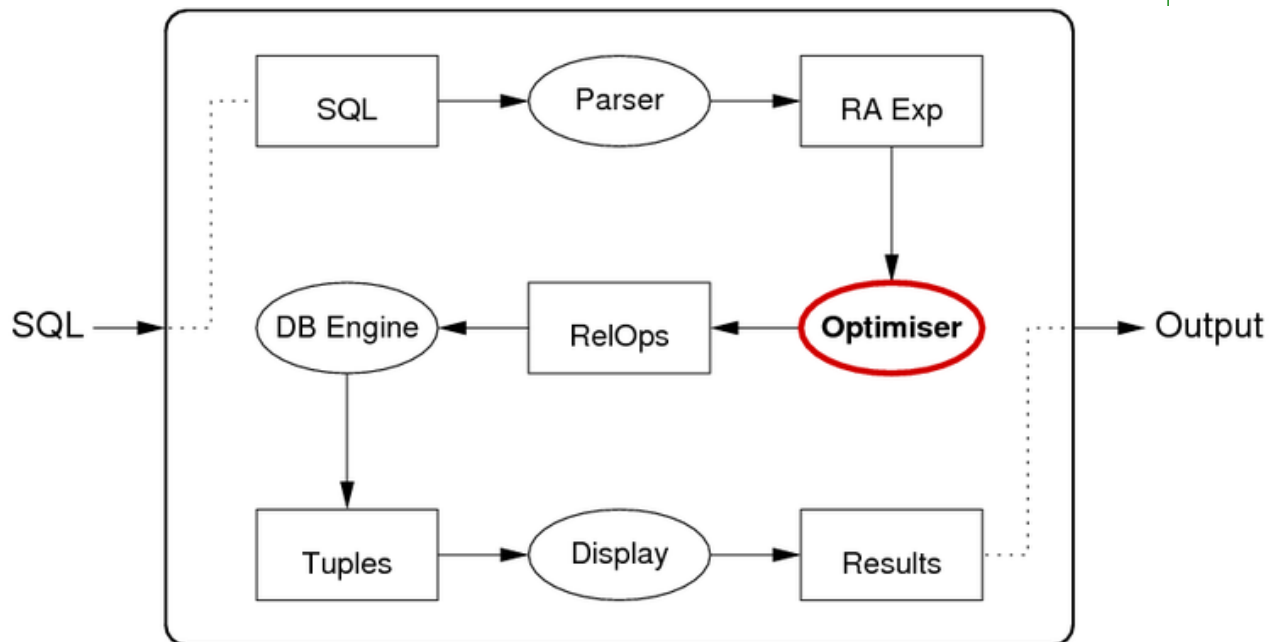
# Query Optimisation

---

- Query Optimisation
- Approaches to Optimisation
- Cost-based Query Optimiser
- Cost Models and Analysis
- Choosing Access Methods (RelOps)
- PostgreSQL Query Optimization

## ❖ Query Optimisation

Query optimiser: RA expression  $\rightarrow$  efficient evaluation plan



## ❖ Query Optimisation (cont)

Query optimisation is a critical step in query evaluation.

The query optimiser

- takes relational algebra expression from SQL compiler
- produces sequence of RelOps to evaluate the expression
- query execution plan should provide efficient evaluation

"Optimisation" is a misnomer since query optimisers

- aim to find a good plan ... but maybe not optimal

Observed Query Time = Planning time + Evaluation time

## ❖ Query Optimisation (cont)

---

Why do we not generate optimal query execution plans?

Finding an optimal query plan ...

- requires exhaustive search of a **space of possible plans**
- for each possible plan, need to estimate cost (not cheap)

Even for relatively small query, search space is *very large*.

Compromise:

- do limited search of query plan space (guided by heuristics)
- *quickly* choose a *reasonably efficient* execution plan

## ❖ Approaches to Optimisation

---

Three main classes of techniques developed:

- algebraic    (equivalences, rewriting, heuristics)
- physical    (execution costs, search-based)
- semantic    (application properties, heuristics)


All driven by aim of minimising (or at least reducing) "cost".

Real query optimisers use a combination of algebraic+physical.

Semantic QO is good idea, but expensive/difficult to implement.

## ❖ Approaches to Optimisation (cont)

Example of optimisation transformations:



For join, may also consider sort/merge join and hash join.

## ❖ Cost-based Query Optimiser

Approximate algorithm for cost-based optimisation:

```

translate SQL query to RAexp
for enough transformations RA' of RAexp {
    while (more choices for RelOps) {
        Plan = {};  i = 0;  cost = 0
        for each node e of RA' (recursively) {
            ROp = select RelOp method for e
            Plan = Plan U ROp
            cost += Cost(ROp) // using child info
        }
        if (cost < MinCost)
            { MinCost = cost;  BestPlan = Plan }
    }
}

```

Heuristics: push selections down, consider only left-deep join trees.

## ❖ Cost Models and Analysis

---

The cost of evaluating a query is determined by:

- size of relations (database relations and temporary relations)
- access mechanisms (indexing, hashing, sorting, join algorithms)
- size/number of main memory buffers (and replacement strategy)

Analysis of costs involves *estimating*:

- size of intermediate results
- number of disk reads/writes



## ❖ Choosing Access Methods (RelOps)

Performed for each node in RA expression tree ...

Inputs:

- a single RA operation ( $\sigma$ ,  $\pi$ ,  $\bowtie$ )
- information about file organisation, data distribution, ...
- list of operations available in the database engine

Output:

- specific DBMS operation to implement this RA operation

## ❖ Choosing Access Methods (RelOps) (cont)

### Example:

- RA operation:  $Sel_{[name='John' \wedge age > 21]}(Student)$
- **Student** relation has B-tree index on **name**
- database engine (obviously) has B-tree search method

giving

```
tmp[i]      := BtreeSearch[name='John'](Student)
tmp[i+1]    := LinearSearch[age>21](tmp[i])
```

Where possible, use pipelining to avoid storing **tmp[i]** on disk.

## ❖ Choosing Access Methods (RelOps) (cont)

Rules for choosing  $\sigma$  access methods:

- $\sigma_{A=c}(R)$  and  $R$  has index on  $A \Rightarrow$   
**indexSearch[A=c](R)**
- $\sigma_{A=c}(R)$  and  $R$  is hashed on  $A \Rightarrow$  **hashSearch[A=c](R)**
- $\sigma_{A=c}(R)$  and  $R$  is sorted on  $A \Rightarrow$   
**binarySearch[A=c](R)**
- $\sigma_{A \geq c}(R)$  and  $R$  has clustered index on  $A$   
 $\Rightarrow$  **indexSearch[A=c](R)** then scan
- $\sigma_{A \geq c}(R)$  and  $R$  is hashed on  $A$   
 $\Rightarrow$  **linearSearch[A>=c](R)**

## ❖ Choosing Access Methods (RelOps) (cont)

Rules for choosing  $\bowtie$  access methods:

- $R \bowtie S$  and  $\mathbf{R}$  fits in memory buffers  $\Rightarrow$  **bnlJoin(R, S)**
- $R \bowtie S$  and  $\mathbf{S}$  fits in memory buffers  $\Rightarrow$  **bnlJoin(S, R)**
- $R \bowtie S$  and  $\mathbf{R, S}$  sorted on join attr  $\Rightarrow$  **smJoin(R, S)**
- $R \bowtie S$  and  $\mathbf{R}$  has index on join attr  $\Rightarrow$  **inlJoin(S, R)**
- $R \bowtie S$  and no indexes, no sorting  $\Rightarrow$  **hashJoin(R, S)**

(**bnl** = block nested loop; **inl** = index nested loop; **sm** = sort merge)

## ❖ PostgreSQL Query Optimization

---

Input: tree of **Query** nodes returned by parser

Output: tree of **Plan** nodes used by query **executor**

- wrapped in a **PlannedStmt** node containing state info

Intermediate data structures are trees of **Path** nodes

- a path tree represents one evaluation order for a query

All **Node** types are defined in **include/nodes/\*.h**

## ❖ PostgreSQL Query Optimization (cont)

Query optimisation proceeds in two stages (after parsing)...

### Rewriting:

- uses PostgreSQL's **rule** system
- query tree is expanded to include e.g. view definitions

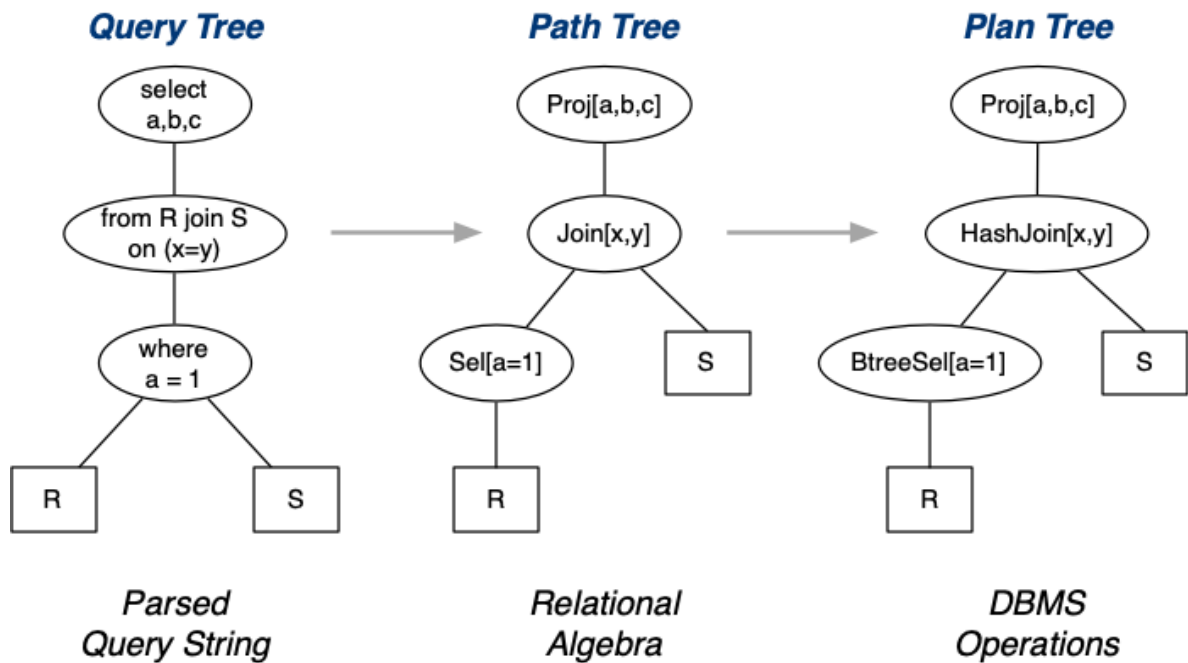
### Planning and optimisation:

- using cost-based analysis of generated paths
- via one of **two** different path generators
- chooses least-cost path from all those considered

Then produces a **Plan** tree from the selected path.

## ❖ PostgreSQL Query Optimization (cont)

select a,b,c from R join S on (x=y) where a = 1



Produced: 5 Apr 2021