

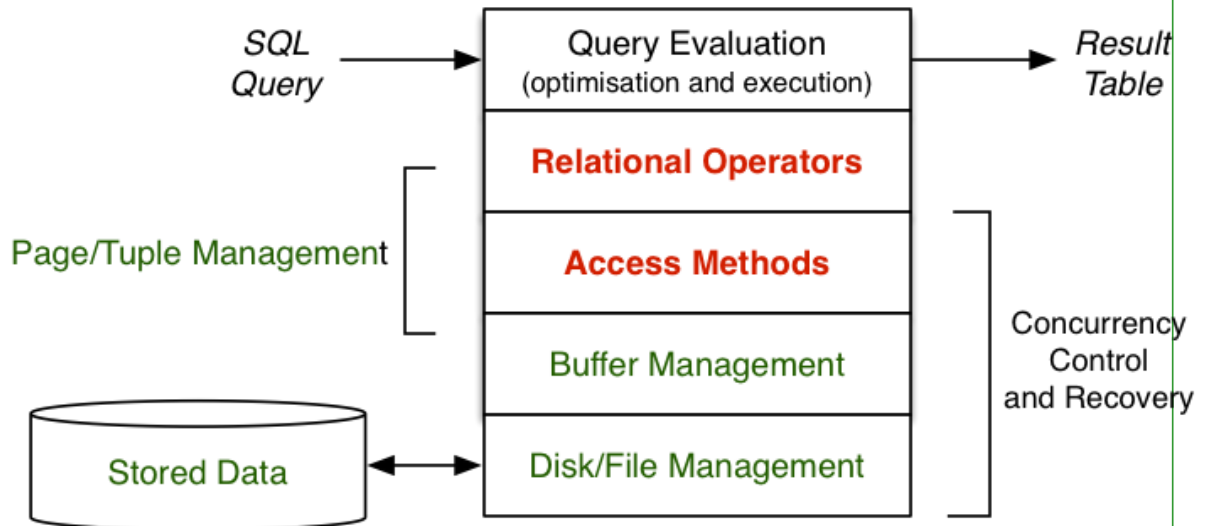
# Relational Operations

---

- DBMS Architecture (revisited)
- Relational Operations
- Cost Models
- Query Types

## ❖ DBMS Architecture (revisited)

Implementation of relational operations in DBMS:



COMP9315 21T1 ◇ Relational Operations ◇ [1/11]

## ❖ Relational Operations

DBMS core = relational engine, with implementations of

- selection, projection, join, set operations
- scanning, sorting, grouping, aggregation, ...

In this part of the course:

- examine methods for implementing each operation
- develop cost models for each implementation
- characterise when each method is most effective

Terminology reminder:

- tuple = collection of data values under some schema  $\cong$  record
- page = block = collection of tuples + management data = i/o unit
- relation = table  $\cong$  file = collection of tuples

## ❖ Relational Operations (cont)

In order to implement relational operations the low-levels of the system provides:

- **Relation** `openRel(db,name)`
  - get handle on relation **name** in database **db**
- **Page** `request_page(rel,pid)`
  - get page **pid** from relation **rel**, return buffer containing page
- **Record** `get_record(buf,tid)`
  - return record **tid** from page **buf**
- **Tuple** `mkTuple(rel,rec)`
  - convert record **rec** to a tuple, based on **rel** schema

## ❖ Relational Operations (cont)

### Example of using low-level functions

```
// scan a relation Emps
Page p; // current page
Tuple t; // current tuple
Relation r = relOpen(db, "Emps");
for (int i = 0; i < nPages(r); i++) {
    p = request_page(rel, i);
    for (int j = 0; j < nRecs(p); j++)
        t = mkTuple(r, get_record(p, j));
        ... process tuple t ...
    }
}
```

## ❖ Relational Operations (cont)

Two "dimensions of variation":

- which relational operation (e.g. Sel, Proj, Join, Sort, ...)
- which access-method (e.g. file struct: heap, indexed, hashed, ...)

Each **query method** involves an operator and a file structure:

- e.g. primary-key selection on hashed file
- e.g. primary-key selection on indexed file
- e.g. join on ordered heap files (sort-merge join)
- e.g. join on hashed files (hash join)
- e.g. two-dimensional range query on R-tree indexed file

We are interested in *cost* of query methods (and insert/delete operations)

## ❖ Relational Operations (cont)

### SQL vs DBMS engine

- **select ... from R where C**
  - find relevant tuples (satisfying C) in file(s) of R
- **insert into R values(...)**
  - place new tuple in some page of a file of R
- **delete from R where C**
  - find relevant tuples and "remove" from file(s) of R
- **update R set ... where C**
  - find relevant tuples in file(s) of R and "change" them

## ❖ Cost Models

---

An important aspect of this course is

- analysis of cost of various query methods

**Cost** can be measured in terms of

- *Time Cost*: total time taken to execute method, or
- *Page Cost*: number of pages read and/or written

Primary assumptions in our cost models:

- memory (RAM) is "small", fast, byte-at-a-time
- disk storage is very large, slow, page-at-a-time



## ❖ Cost Models (cont)

Since *time cost* is affected by many factors

- speed of i/o devices (fast/slow disk, SSD)
- load on machine

we do not consider time cost in our analyses.

For comparing methods, *page cost* is better

- identifies workload imposed by method
- BUT is clearly affected by buffering


Estimating costs with multiple concurrent ops *and* buffering is difficult!!

Additional assumption: every page request leads to some i/o

## ❖ Cost Models (cont)

In developing cost models, we also assume:

- a relation is a set of  $r$  tuples, with average tuple size  $R$  bytes
- the tuples are stored in  $b$  data pages on disk
- each page has size  $B$  bytes and contains up to  $c$  tuples
- the tuples which answer query  $q$  are contained in  $b_q$  pages
- data is transferred disk  $\leftrightarrow$  memory in whole pages
- cost of disk  $\leftrightarrow$  memory transfer  $T_{r/w}$  is very high



## ❖ Cost Models (cont)

Our cost models are "rough" (based on assumptions)

But do give an  $O(x)$  feel for how expensive operations are.

Example "rough" estimation: how many piano tuners in Sydney?

- Sydney has  $\approx 4\,000\,000$  people
- Average household size  $\approx 3 \therefore 1\,300\,000$  households
- Let's say that 1 in 10 households owns a piano
- Therefore there are  $\approx 130\,000$  pianos
- Say people get their piano tuned every 2 years (on average)
- Say a tuner can do 2/day, 250 working-days/year
- Therefore 1 tuner can do 500 pianos per year
- Therefore Sydney would need  $\approx 130000/2/500 = 130$  tuners

Actual number of tuners in Yellow Pages = 120

Example borrowed from Alan Fekete at Sydney University.

## ❖ Query Types

Type	SQL	RelAlg	a.k.a.
Scan	<b>select * from R</b>	$R$	-
Proj	<b>select <math>x,y</math> from R</b>	$Proj[x,y]R$	-
Sort	<b>select * from R order by <math>x</math></b>	$Sort[x]R$	<i>ord</i>
Sel <sub>1</sub>	<b>select * from R where id = <math>k</math></b>	$Sel[id=k]R$	<i>one</i>
Sel <sub>n</sub>	<b>select * from R where a = <math>k</math></b>	$Sel[a=k]R$	-
Join <sub>1</sub>	<b>select * from R,S where R.id = S.r</b>	$R Join[id=r] S$	-

Different query classes exhibit different query processing behaviours.

Produced: 27 Feb 2021