

Sorted Files

- Sorted Files
- Selection in Sorted Files
- Insertion into Sorted Files
- Deletion from Sorted Files

COMP9315 21T1 ◇ Sorted Files ◇ [0/12]

❖ Sorted Files


Records stored in file in order of some field **k** (the sort key).

Makes searching more efficient; makes insertion less efficient

E.g. assume $c = 4$


a,b,c,d	e,f,h,i	j,k,l,m
---------	---------	---------

insert g



❖ Sorted Files (cont)

In order to mitigate insertion costs, use overflow pages.



Total number of overflow pages = b_{ov} .

Average overflow chain length = $ov = b_{ov}/b$.

Bucket = data page + its overflow page(s)

❖ Selection in Sorted Files

For *one* queries on sort key, use binary search.

```
// select * from R where k = val  (sorted on R.k)
lo = 0; hi = nPages(rel)-1
while (lo <= hi) {
    mid = (lo+hi) / 2;  // int division with truncation
    (tup, loVal, hiVal) = searchBucket(rel, mid, x, val);
    if (tup != NULL) return tup;
    else if (val < loVal) hi = mid - 1;
    else if (val > hiVal) lo = mid + 1;
    else return NOT_FOUND;
}
return NOT_FOUND;
```

where **rel** is relation handle, **mid**, **lo**, **hi** are page indexes,
k is a field/attr, **val**, **loVal**, **hiVal** are values for **k**

❖ Selection in Sorted Files (cont)

Search a page and its overflow chain for a key value

```
searchBucket(rel,p,k,val)
{
    get_page(rel,p,buf);
    (tup,min,max) = searchPage(buf,k,val,+INF,-INF)
    if (tup != NULL) return(tup,min,max);
    ovf = openOvFile(f);
    ovp = overflow(buf);
    while (tup == NULL && ovp != NO_PAGE) {
        get_page(ovf,ovp,buf);
        (tup,min,max) = searchPage(buf,k,val,min,max)
        ovp = overflow(buf);
    }
    return (tup,min,max);
}
```

Assumes each page contains index of next page in Ov chain

❖ Selection in Sorted Files (cont)

Search within a page for key; also find min/max key values

```
searchPage(buf, k, val, min, max)
{
    res = NULL;
    for (i = 0; i < nTuples(buf); i++) {
        tup = get_tuple(buf, i);
        if (tup.k == val) res = tup;
        if (tup.k < min) min = tup.k;
        if (tup.k > max) max = tup.k;
    }
    return (res, min, max);
}
```

❖ Selection in Sorted Files (cont)

The above method treats each bucket like a single large page.

Cases:

- best: find tuple in first data page we read
- worst: full binary search, and not found
 - examine $\log_2 b$ data pages
 - plus examine all of their overflow pages
- average: examine some data pages + their overflow pages

$Cost_{one}$: Best = 1 Worst = $\log_2 b + b_{ov}$


Average case cost analysis needs assumptions (e.g. data distribution)

❖ Selection in Sorted Files (cont)

For *pmr* query, on non-unique attribute k , where file is sorted on k

- tuples containing k may span several pages

E.g. **select * from R where k = 2**



Begin by locating a page p containing $k=val$ (as for *one* query).

Scan backwards and forwards from p to find matches.


Thus, $Cost_{pmr} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$

❖ Selection in Sorted Files (cont)

For *range* queries on unique sort key (e.g. primary key):

- use binary search to find lower bound
- read sequentially until reach upper bound

E.g. **select * from R where k >= 5 and k <= 13**




$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

❖ Selection in Sorted Files (cont)

For *range* queries on non-unique sort key, similar method to *pmr*.

- binary search to find lower bound
- then go backwards to start of run
- then go forwards to last occurrence of upper-bound

E.g. **select * from R where k >= 2 and k <= 6**



$$Cost_{range} = Cost_{one} + (b_q - 1) \cdot (1 + Ov)$$

❖ Selection in Sorted Files (cont)

So far, have assumed query condition involves sort key k .

But what about **`select * from R where j = 100.0?`**

If condition contains attribute j , not the sort key

- file is unlikely to be sorted by j as well
- sortedness gives no searching benefits

$Cost_{one}$, $Cost_{range}$, $Cost_{pmr}$ as for heap files

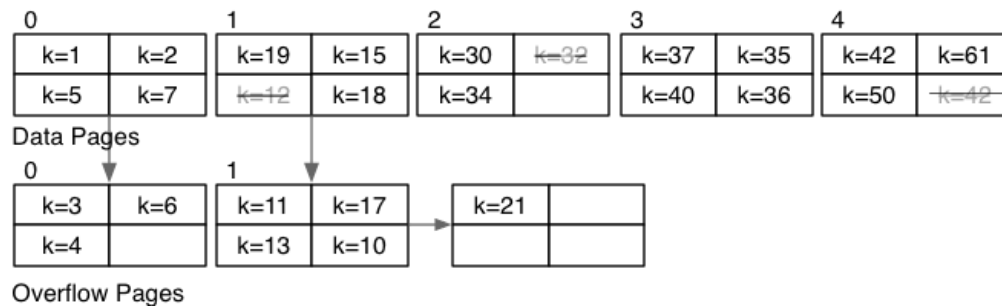
❖ Insertion into Sorted Files

Insertion approach:

- find appropriate page for tuple (via binary search)
- if page not full, insert into page
- otherwise, insert into next overflow page with space

Thus, $Cost_{insert} = Cost_{one} + \delta_w$ (where $\delta_w = 1$ or 2)

Consider insertions of $k=33$, $k=25$, $k=99$ into:



❖ Deletion from Sorted Files

E.g. **delete from R where k = 2**

Deletion strategy:

- find matching tuple(s)
- mark them as deleted

Cost depends on **selectivity** of selection condition

Recall: selectivity determines b_q (# pages with matches)

Thus, $Cost_{delete} = Cost_{select} + b_{qw}$

Produced: 7 Mar 2021