

Week 4 Exercises

- Things to Note
- Assignment 1
- DBMS Architecture (revisited)
- Exercise: File Merging
- Projection
- Sort-based Projection
- Exercise: Sort-based Projection
- Hash-based Projection
- Exercise: Hash-based Projection
- Query Types
- Exercise: Query Types
- Data File Structures
- Exercise: Cost of Deletion in Heaps
- Exercise: Searching in Sorted File
- Exercise: Optimising Sorted-file Search
- Exercise: Insertion into Static Hashed File
- Things to Note
- Hash Values and Bit-strings
- Exercise: Bit Manipulation
- Hashing
- Hashing Performance
- Hashed Files
- Linear-hashed Files
- Exercise: Insertion into Linear Hashed File

COMP9315 21T1 ◇ Week 4 Exercises ◇ [0/29]

❖ Things to Note

Quiz 2 ... due Friday 11 March at 9pm

Assignment 1 due Friday 18 March at 9pm

"Disk quota exceeded"? Check your disk usage on /localstorage/

- use the **du** command to find "hot spots"
- if log is big: stop server, remove log file, restart server

This weekend is Census Weekend ...

- if you want to drop a course and not pay for it, drop it now
- if you want to drop COMP9315, drop it in MyUNSW and email me

❖ Assignment 1

Debugging the assignment ...

- can't easily use GDB or vscode; need debugging print's
- can use **ereport()** and **elog()**; see PG Docs 54.2

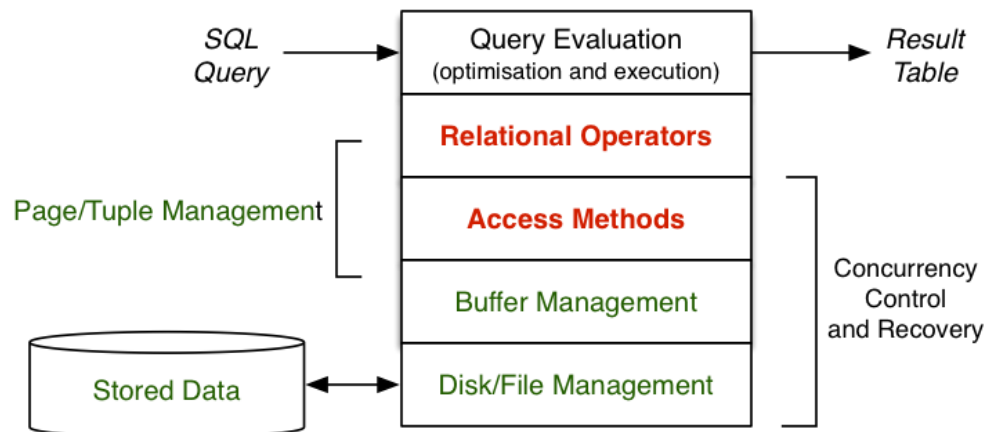
Implementing **PersonName** ...

- three data usages: readable, computable, storeable
 - readable → storeable in **pname_in()**
 - storeable → readable in **pname_out()**
 - storeable → computable in e.g. **pname_cmp()**
- reminder: pointers to **malloc**'d memory structures are not storeable

Rebuilding PostgreSQL will *not* solve problems in **pname.***

❖ DBMS Architecture (revisited)

Implementation of relational operations in DBMS:



❖ Exercise: File Merging

Implement a merging algorithm

- for two sorted files, using 3 buffers, with $b_1=5$, $b_2=3$
- for one unsorted file, using 3 buffers, with $b = 12$
- for one unsorted file, using 5 buffers, with $b = 27$

Assume that we have functions

- **get_page(rel, pid, buf)** ... read specified page into buffer
- **put_page(rel, pid, buf)** ... write a page to disk, at position pid
- **clear_page(rel, buf)** ... make page have zero tuples
- **sort_page(buf)** ... in-memory sort of tuples in page
- **nPages(rel), nTuples(buf), get_tuple(buf, tid)**

❖ Projection

Projection removes some attributes from tuples

Projection can produce duplicates

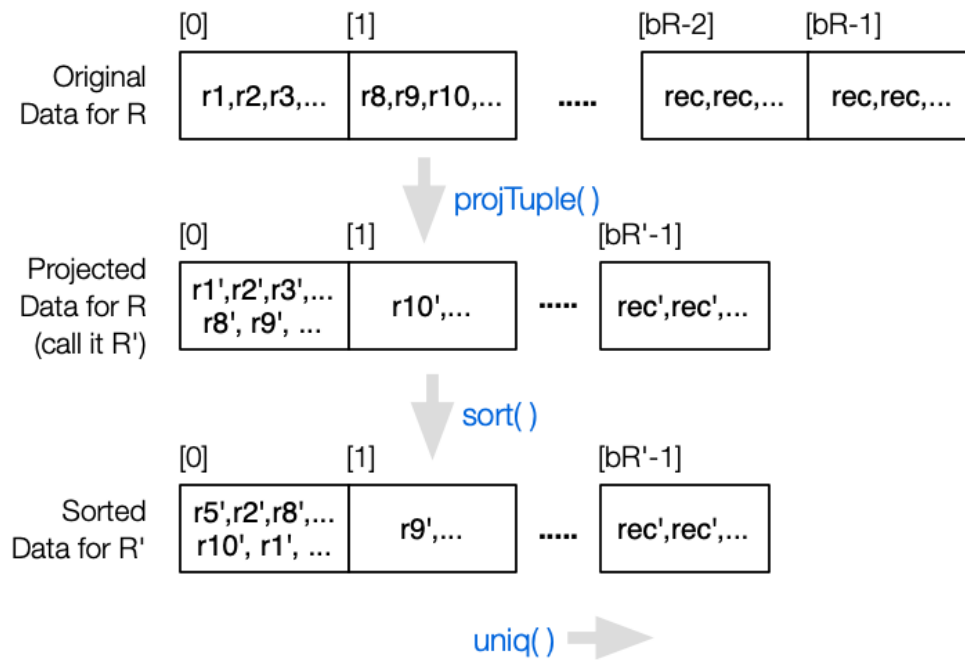
$R(x,y,z) = [(a,b,6), (a,c,7), (a,b,8)]$

$\text{Proj}[x,y] R \Rightarrow [(a,b), (a,c), (a,b)]$

If duplicates are ok, projection is trivial

```
for each tuple t {  
    t' = t projected on some attributes  
    output t'  
}
```

❖ Sort-based Projection



❖ Sort-based Projection (cont)

Duplicate removal in a sorted file

```
curr = NULL
for each tuple t {
    if (t == curr) continue
    output t
    curr = t
}
```

Requires scan of all b_{in} pages in input; writes b_{out} pages; $b_{out} \leq b_{in}$

❖ Exercise: Sort-based Projection

Consider a table $R(x,y,z)$ with tuples:

Page 0:	(1,1,'a')	(11,2,'a')	(3,3,'c')
Page 1:	(13,5,'c')	(2,6,'b')	(9,4,'a')
Page 2:	(6,2,'a')	(17,7,'a')	(7,3,'b')
Page 3:	(14,6,'a')	(8,4,'c')	(5,2,'b')
Page 4:	(10,1,'b')	(15,5,'b')	(12,6,'b')
Page 5:	(4,2,'a')	(16,9,'c')	(18,8,'c')

SQL: **create T as (select distinct y from R)**


Assuming:

- 3 memory buffers, 2 for input, one for output
- pages/buffers hold 3 **R** tuples (i.e. $c_R=3$), 6 **T** tuples (i.e. $c_T=6$)

Show how sort-based projection would execute this statement.


❖ Hash-based Projection

Partitioning phase:



❖ Hash-based Projection (cont)

Duplicate elimination phase:



❖ Hash-based Projection (cont)

Pseudo-code for hash-based projection (with N buffers)

```
open N-1 partition files (f[0], f[1], ... f[N-2])
clear all buffers (b[0], b[1], ... b[N-2])
for each tuple t {    -- via input buffer
    h = h1(t)
    add t to b[h]
    if (b[h] is full) { write to f[h]; clear b[h] }
}
rewind all files f[i]
for each partition p in 0 .. N-1 {
    for each tuple t in f[p] {    -- via input buffer
        h = h2(t)
        if (t in b[h]) continue
        add t to b[h]
    }
    for each buffer b[i] write b[i] to output file
}
```

❖ Exercise: Hash-based Projection

Consider a table $R(x,y,z)$ with tuples:

```

Page 0:  (1,1,'a')   (11,2,'a')   (3,3,'c')
Page 1:  (13,5,'c')  (2,6,'b')   (9,4,'a')
Page 2:  (6,2,'a')   (17,7,'a')  (7,3,'b')
Page 3:  (14,6,'a')  (8,4,'c')   (5,2,'b')
Page 4:  (10,1,'b')  (15,5,'b')  (12,6,'b')
Page 5:  (4,2,'a')   (16,9,'c')  (18,8,'c')

```

-- and then the same tuples repeated for pages 6-11

SQL: **create T as (select distinct y from R)**

Assuming:

- 4 memory buffers, one for input, 3 for partitioning
- pages/buffers hold 3 **R** tuples (i.e. $c_R=3$), 4 **T** tuples (i.e. $c_T=4$)
- hash functions: $h_1(x) = x\%3$, $h_2(x) = (x\%4)\%3$

Show how hash-based projection would execute this statement.

❖ Query Types

Type	SQL	RelAlg	a.k.a.
Scan	select * from R	R	-
Proj	select x,y from R	$Proj[x,y]R$	-
Sort	select * from R order by x	$Sort[x]R$	<i>ord</i>
Sel_1	select * from R where id = k	$Sel[id=k]R$	<i>one</i>
Sel_n	select * from R where a = k	$Sel[a=k]R$	-
$Join_1$	select * from R,S where R.id = S.r	$R Join[id=r] S$	-

Different query classes exhibit different query processing behaviours.

❖ Exercise: Query Types

Using the relation:

```
create table Courses (  
  id          integer primary key,  
  code        char(8),    -- e.g. 'COMP9315'  
  title       text,       -- e.g. 'Computing 1'  
  year        integer,    -- e.g. 2000..2016  
  convenor    integer references Staff(id),  
  constraint once_per_year unique (code,year)  
);
```

give examples of each of the following query types:

1. a 1-d *one* query, an n-d *one* query
2. a 1-d *pmr* query, an n-d *pmr* query
3. a 1-d *range* query, an n-d *range* query

Suggest how many solutions each might produce ...

❖ Data File Structures

Three common file types ($b = 5, c = 4$)

	[0]				[4]
Heap File	k=1, k=14, k=7, k=10	k=2, k=6, k=3, k=15	k=13, k=4, k=20, k=12	k=19, k=11, k=7, k=16	k=17, k=8, k=18, k=5

	[0]				[4]
Sorted File	k=1, k=2, k=3, k=4	k=5, k=6, k=7, k=8	k=9, k=10, k=11, k=12	k=13, k=14, k=15, k=16	k=17, k=18, k=19, k=20

	[0]				[4]
Hash File $h(x) = x \% 5$	k=5, k=10, k=15, k=20	k=1, k=6, k=11, k=16	k=2, k=7, k=12, k=17	k=3, k=8, k=13, k=18	k=4, k=9, k=14, k=19

❖ Exercise: Cost of Deletion in Heaps

Consider the following queries ...

```
delete from Employees where id = 12345    -- one
delete from Employees where dept = 'Marketing'  -- pmr
delete from Employees where 40 <= age and age < 50  -- range
```

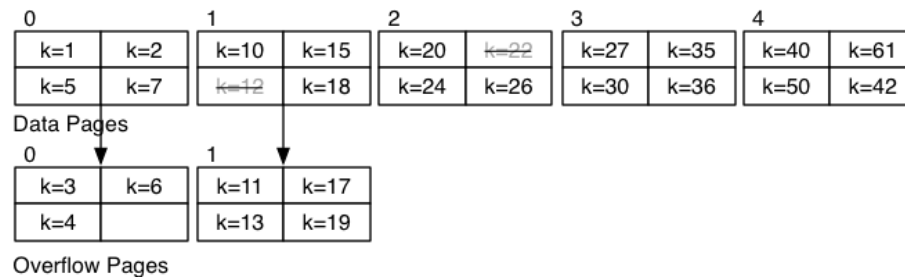
Show how each will be executed and estimate the cost, assuming:

- $b = 100$, $b_{q2} = 3$, $b_{q3} = 20$

State any other assumptions.

❖ Exercise: Searching in Sorted File

Consider this sorted file with overflows ($b=5, c=4$):



Compute the cost for answering each of the following:

- **select * from R where k = 24**
- **select * from R where k = 3**
- **select * from R where k = 14**
- **select max(k) from R**

❖ Exercise: Optimising Sorted-file Search

The **searchBucket (f, p, k, val)** function requires:

- read the p^{th} page from data file
- scan it to find a match and min/max k values in page
- while no match, repeat the above for each overflow page
- if we find a match in any page, return it
- otherwise, remember min/max over all pages in bucket

Suggest an optimisation that would improve **searchBucket ()** performance for most buckets.

❖ Exercise: Insertion into Static Hashed File

Consider a file with $b=4$, $c=3$, $d=2$, $h(x) = \text{bits}(d, \text{hash}(x))$

Insert tuples in alpha order with the following keys and hashes:

k	$\text{hash}(k)$	k	$\text{hash}(k)$	k	$\text{hash}(k)$	k	$\text{hash}(k)$
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

❖ Things to Note

Quiz 2 ... due Friday 11 March (tomorrow!) at 9pm

Assignment 1 ... due Friday 18 March at 9pm

- need a hash function? write your own, or use a PostgreSQL built-in one

Assignment 2 ... implement a [multi-attribute linear-hashed file](#)

This weekend is Census Weekend ...

- if you want to drop a course and not pay for it, drop it now
- if you want to drop COMP9315, drop it in MyUNSW and email me

❖ Hash Values and Bit-strings

Hashing requires $h(k) :: \text{KeyVal} \rightarrow \text{HashVal}$

HashVal is typically a 32-bit integer, which is mapped to $0 .. b-1$

For arbitrary b , mapping done via **$\text{PageID} = h(k) \% b$**

If $b = 2^d$, mapping can be done via bitwise AND

E.g. $b = 8$, **$\text{PageID} = h(k) \& 0b0111$**

For any d , use a mask with lower-order d bits set to 1

❖ Exercise: Bit Manipulation

1. Write a function to display **uint32** values as **01010110...**

```
char *showBits(uint32 val, char *buf);
```

(assumes supplied buffer is large enough, like **gets()**)


2. Write a function to extract the d bits of a **uint32**

```
uint32 bits(int d, uint32 val);
```

If $d > 0$, gives low-order bits; if $d < 0$, gives high-order bits

❖ Hashing

Basic idea: use key value to compute page address of tuple.



e.g. tuple with key = v is stored in page i

Requires: hash function $h(v)$ that maps $KeyVal \rightarrow [0..b-1]$.

- hashing converts key value (any type) into integer value
- integer value is then mapped to page index
- note: can view integer value as a bit-string

❖ Hashing Performance

Aims:

- distribute tuples evenly amongst buckets (data + overflow pages)
- have most data pages nearly full (minimise wasted space)

Note: if data distribution not uniform, address distribution can't be uniform.

Best case: every bucket contains same number of tuples.

Worst case: every tuple hashes to same bucket.

Average case: some buckets have more tuples than others.

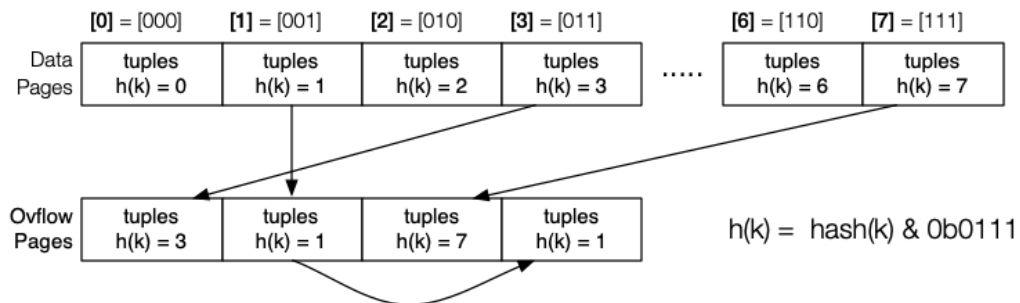
Use overflow pages to handle "overflow" buckets (cf. sorted files)

All tuples in each bucket have same hash value.

❖ Hashed Files

Hash function maps key $\rightarrow 0 \dots b-1$

E.g. if $b = 2^d$, simply use lower-order d bits of hash




Static hashing: fixed-size file \Rightarrow overflow chains grow without bound

Linear hashing: variable-size file, using sp = split pointer, d = depth

❖ Linear-hashed Files

If $b = 2^d$, the file behaves exactly like standard hashing.

If $b \neq 2^d$, treat different parts of the file differently.



Parts *A* and *C* are treated as if part of a file of size 2^{d+1}

Part *B* is treated as if part of a file of size 2^d

Part *D* does not yet exist (tuples in *B* may eventually move into it)

❖ Linear-hashed Files (cont)

Mapping a hash value to a bucket id for searching

```
// select * from R where k = val
h = hash(val);
pid = bits(d,h); // take lower-order d bits from h
if (pid < sp) { pid = bits(d+1,h); }

P = getPage(datafile(reln), pid)
for each tuple t in page P
    if (t.k == val) add t to answerSet

pid = overflow(P)
while (pid != NULL) {
    P = getPage(ovfile(reln), pid)
    for each tuple t in page P
        if (t.k == val) add t to answerSet
    pid = overflow(P)
}
```

❖ Linear-hashed Files (cont)

Periodically, extend file by one page

- new page has address $sp + 2^d$
- tuples in **bucket** sp are redistributed
 - by considering $d+1$ bits of hash (e.g. **0111** → **?0111**)
 - tuples where extra bit is 0, stay in bucket sp
 - tuples where extra bit is 1, move to bucket $sp + 2^d$
- after this, $sp++$; if $(sp == 2^d) \{ sp = 0; d++ \}$

This process is called **splitting**

It can reduce, but may not remove, overflow chain for bucket sp

❖ Exercise: Insertion into Linear Hashed File

Consider a file with $b=4$, $c=3$, $d=2$, $sp=0$, $hash(x)$ as below

Insert tuples in alpha order with the following keys and hashes:

k	$hash(k)$	k	$hash(k)$	k	$hash(k)$	k	$hash(k)$
a	10001	g	00000	m	11001	s	01110
b	11010	h	00000	n	01000	t	10011
c	01111	i	10010	o	00110	u	00010
d	01111	j	10110	p	11101	v	11111
e	01100	k	00101	q	00010	w	10000
f	00010	l	00101	r	00000	x	00111

The hash values are the 5 lower-order bits from the full 32-bit hash.

Split *before* every sixth insert.

Produced: 10 Mar 2022