

Week 5 Exercises

- Things to Note
- Exercise: Index Storage Overheads
- Exercise: One vs Multiple Indices
- Bitmap Indexes
- Exercise: Bitmap Index
- Exercise: Insertion into B-trees
- Exercise: Multi-attribute Hashing
- Exercise: Partial hash values in MAH
- Exercise: Representing Stars
- Exercise: MA.Hashing Query Cost
- Exercise: MA.Hashing Design
- Exercise: Query Types and Tuple Space
- Exercise: Searching in kd-Trees
- Exercise: Searching in Quad-trees
- Exercise: Query with R-trees

❖ Things to Note

- More Exercises now available
- Assignment 1 due before 9pm Friday
- Assignment 2 available next week
- Next week is Flexibility Week
 - no new Videos+Slides
 - no online sessions
 - but Quiz 3 will run

❖ Exercise: Index Storage Overheads

Consider a relation with the following storage parameters:

- $B = 8192$, $R = 128$, $r = 100000$
- header in data pages: 256 bytes
- key is integer, data file is sorted on key
- index entries (keyVal,tupleID): 8 bytes
- header in index pages: 32 bytes

How many pages are needed to hold a dense index?

How many pages are needed to hold a sparse index?

❖ Exercise: One vs Multiple Indices

Consider a relation with $r = 100,000$, $B = 4096$, defined as:

```
create table Students (  
    id          integer primary key,  
    name        char(10), -- simplified  
    gender      char(1),  -- 'm', 'f', '?'  
    birthday    char(5)   -- 'MM-DD'  
);
```

Assumptions:

- data file is not ordered on any attribute
- has a dense B-tree index on each attribute
- 96 bytes of header in each data/index page

❖ Exercise: One vs Multiple Indices (cont)

For **Students (id, name, gender, birthday) ...**

- calculate the size of the data file and each index
- describe the selectivity of each attribute

Now consider a query on this relation:

```
select * from Students
where  name='John' and birthday='04-01'
```

- estimate the cost of answering using **name** index
- estimate the cost of answering using **birthday** index
- estimate the cost of answering using both indices

❖ Bitmap Indexes

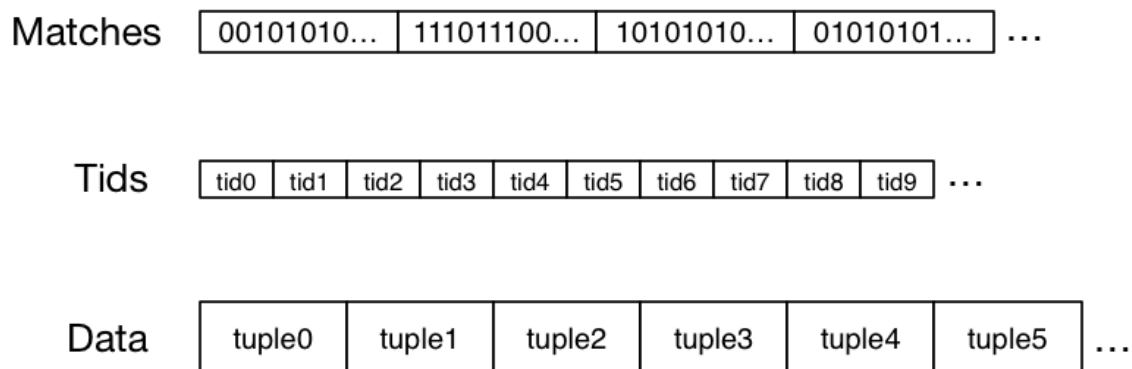
Alternative index structure, focussing on sets of tuples:

Data File				Colour Index	
	Part#	Colour	Price		
[0]	P7	red	\$2.50	red	100011...
[1]	P1	green	\$3.50	blue	001100...
[2]	P9	blue	\$4.10	green	010000...
[3]	P4	blue	\$7.00		
[4]	P5	red	\$5.20		
[5]	P5	red	\$2.50		
.....				Price Index	
				< \$4.00	110001...
				>= \$4.00	001110...

Index contains bit-strings of r bits, one for each value/range

❖ Bitmap Indexes (cont)

Also useful to have a file of **tid**s, one for each tuple:



❖ Bitmap Indexes (cont)

Answering queries using bitmap index:

```
Matches = AllOnes(r)
foreach attribute A with index {
    // select ith bit-string for attribute A
    // based on value associated with A in WHERE
    Matches = Matches & Bitmaps[A][i]
}
// Matches contains 1-bit for each matching tuple
foreach i in 0..r-1 {
    if (Matches[i] == 0) continue;
    Pages = Pages U {pageOf(Tids[i])}
}
foreach pid in Pages {
    P = getPage(pid)
    extract matching tuples from P
}
```


❖ Exercise: Bitmap Index

Using the following file structure:

Data File				Colour Index		Price Index	
	Part#	Colour	Price				
[0]	P7	red	\$2.50	red	100011...	< \$4.00	110001...
[1]	P1	green	\$3.50	blue	001100...		
[2]	P9	blue	\$4.10	green	010000...		
[3]	P4	blue	\$7.00				
[4]	P5	red	\$5.20				
[5]	P5	red	\$2.50				
.....							

Show how the following queries would be answered:

```
select * from Parts
where colour='red' and price < 4.00
```

```
select * from Parts
where colour='green' or colour ='blue'
```

❖ Exercise: Bitmap Index (cont)

Storage costs for bitmap indexes:

- one bitmap for each value/range for each indexed attribute
- each bitmap has length $\text{ceil}(r/8)$ bytes
- e.g. with 50K records and 8KB pages, bitmap fits in one page

Query execution costs for bitmap indexes:

- read one bitmap for each indexed attribute in query
- perform bitwise AND on bitmaps (in memory)
- read pages containing matching tuples

Note: bitmaps could index pages rather than tuples (shorter bitmaps)

❖ Exercise: Insertion into B-trees

Assumptions

- B-tree is initially empty
- each node in the B-tree is a page with $c_i = 5$
- key values are unique (i.e. primary key)
- each leaf node has a pointer to its right sibling
- tuples are added to the data file in order of insertion

Show how the tree grows as we add the following key values

- 100, 50, 80, 200, 20, 65, 150, 120, 110, 75, 10, 180, ...

❖ Exercise: Multi-attribute Hashing

Compute the hash value for the tuple

`('John Smith', 'BSc(CompSci)', 1990, 99.5)`

where $d=6$, $d_1=3$, $d_2=2$, $d_3=1$, and

- $cv = \langle (1,0), (1,1), (2,0), (3,0), (1,2), (2,1), (3,1), (1,3), \dots \rangle$
- $hash_1(\text{'John Smith'}) = \dots 0101010110110100$
- $hash_2(\text{'BSc(CompSci)'}) = \dots 1011111101101111$
- $hash_3(1990) = \dots 0001001011000000$

❖ Exercise: Partial hash values in MAH

Given the following:

- $d=6$, $b=2^6$, $CV = \langle (0,0), (0,1), (1,0), (2,0), (1,1), (0,2), \dots \rangle$
- $\text{hash}(a) = \dots 00101101001101$
- $\text{hash}(b) = \dots 00101101001101$
- $\text{hash}(c) = \dots 00101101001101$

What are the query hashes for each of the following:

- (a,b,c) , $(?,b,c)$, $(a,?,?)$, $(?,?,?)$

❖ Exercise: Representing Stars

Our hash values are bit-strings (e.g. **100101110101**)

MA.Hashing introduces a third value (* = unknown)

How could we represent "bit"-strings like **1011*1*0**010**?

❖ Exercise: MA.Hashing Query Cost

Consider $R(x, y, z)$ using multi-attribute hashing where

$$d = 9 \quad d_x = 5 \quad d_y = 3 \quad d_z = 1$$

How many buckets are accessed in answering each query?

1. **select * from R where x = 4 and y = 2 and z = 1**
2. **select * from R where x = 5 and y = 3**
3. **select * from R where y = 99**
4. **select * from R where z = 23**
5. **select * from R where x > 5**

❖ Exercise: MA.Hashing Design

Consider relation **Person**(**name**, **gender**, **age**) ...

p_Q Query Type Q

0.6 **select age,gender from Person**
 where name=X

0.3 **select count(*) from Person**
 where gender=X

0.1 **select name from Person**
 where gender=X and age=Y

Assume that all other query types have $p_Q=0$.

Design a choice vector to minimise average selection cost.

❖ Exercise: Query Types and Tuple Space

Which part of the tuple-space does each query represent?

Q1: `select * from Rel where X = 'd' and Y = 4`

Q2: `select * from Rel where 'j' < X ≤ 'r'`

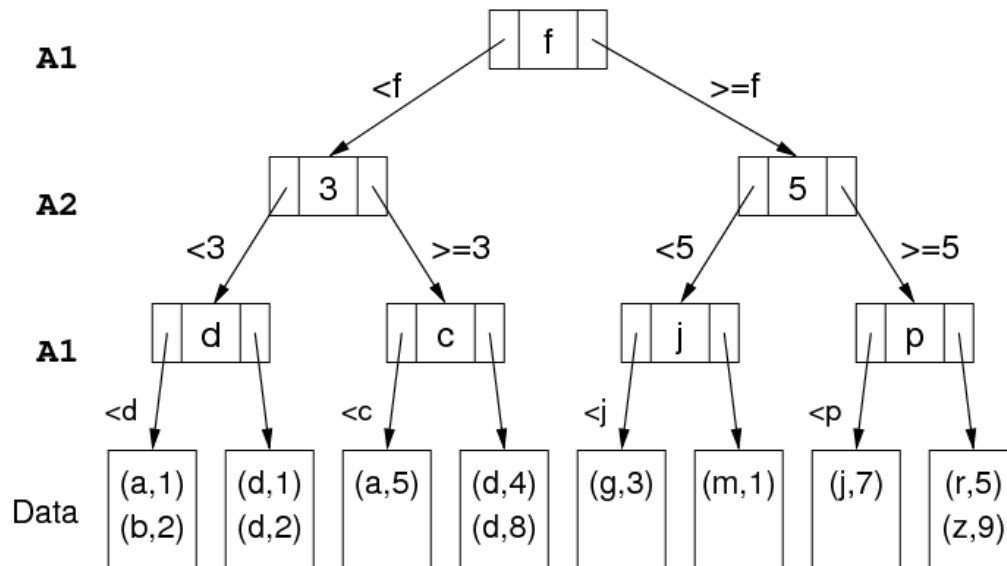
Q3: `select * from Rel where X > 'm' and Y > 4`

Q4: `select * from Rel where 'k' ≤ X ≤ 'p' and 3 ≤ Y ≤ 6`



❖ Exercise: Searching in kd-Trees


Using the following kd-tree index



Answer the queries **(m, 1)**, **(a, ?)**, **(?, 1)**, **(?, ?)**

❖ Exercise: Searching in Quad-trees


Using the following quad-tree index



Answer the queries **(m,1)**, **(a,?)**, **(?,1)**, **(?,?)**

❖ Exercise: Query with R-trees

Using the following R-tree:



Show how the following queries would be answered:

- Q1: `select * from Rel where X='a' and Y=4`
 Q2: `select * from Rel where X='i' and Y=6`
 Q3: `select * from Rel where 'c' ≤ X ≤ 'j' and Y=5`
 Q4: `select * from Rel where X='c'`

Note: can view unknown value $\mathbf{x}=?$ as range $\min(\mathbf{x}) \leq \mathbf{x} \leq \max(\mathbf{x})$

Produced: 16 Mar 2021